

SARAH JENKINS

sarah.jenkins.ai@email.com | +1 555 019 8472 | linkedin.com/in/sarah-jenkins-ai

PROFESSIONAL SUMMARY

Principal-level AI and data science leader with over two decades of experience architecting distributed, enterprise-grade systems and 3+ years building large-scale AI/ML solutions deployed in production. Proven expertise developing predictive models and anomaly-detection workflows, from an autonomous market-analysis agent generating predictive models to evaluation frameworks that detect hallucination and accuracy anomalies in LLM outputs prior to deployment. Deep hands-on skills in Python, JavaScript, and big data technologies (Hadoop, MapReduce, Tableau, Power BI), with a demonstrated record of managing massive, complex datasets - migrating 12,000+ enterprise accounts and roughly 65 million file assets - under strict data security and governance standards. Recognized mentor and subject matter expert who communicates technical findings clearly to both C-suite stakeholders and non-technical audiences.

SKILLS

Python - JavaScript (Node.js) - PyTorch - TensorFlow - Predictive Modeling - Tableau - Power BI - Data Visualization - Big Data Analytics - Hadoop - MapReduce - Data Pipelines - Vector Databases - Enterprise RAG Pipelines - Data Governance - Data Security - Multi-Agent Orchestration & Coordination - LangGraph - Agent-to-Agent (A2A) Protocols - Stateful AI Memory Architectures (Short-Term Context, Long-Term User Memory) - Advanced Prompt Engineering - Model Fine-Tuning & Evaluation - CrewAI - Semantic Kernel - Interruptible Voice Agents - Conversational AI Frameworks - TypeScript - Java - Microservices Architecture - API-First Design - AWS (Bedrock, SageMaker) - Azure AI Studio - GCP Vertex AI - RESTful Services - MuleSoft

EXPERIENCE

Principal AI Architect - NexaCloud Solutions

January 2021 - Present

- Engineered and deployed a multi-agent collaboration framework featuring specialized validation guardrails and an 'LLM-as-judge' agent setup, improving initial output accuracy from 89% to 99% through iterative tiering and self-correction loops.
- Built an LLM-powered marketing automation suite for automated content generation, dynamic lead scoring, and intelligent classification and escalation of inbound customer communications.
- Architected a modular AI ecosystem deploying specialized agents as microservices, utilizing a structured Agent-to-Agent (A2A) protocol alongside a custom MCP server for secure tool-calling into legacy enterprise platforms while safeguarding sensitive enterprise data.
- Designed stateful AI workflows using LangGraph, implementing robust memory management that preserves real-time conversational context and long-term customer preference profiles.
- Led the technical assessment and strategic roadmap for modernizing a core enterprise policy platform, delivering over 450 functional enhancements via strict release governance and change management.

Senior Enterprise Consultant - Vertex Consulting Group

August 2016 - December 2020

- Transformed over 250 legacy reporting structures into optimized, real-time Tableau and Power BI dashboard visualization suites, applying advanced UX principles to drastically improve operational visibility for C-suite stakeholders.
- Spearheaded modernization initiatives migrating legacy mainframe applications to modern cloud architectures, orchestrating data pipelines that successfully migrated over 12,000 enterprise accounts and associated records while maintaining data integrity.
- Directed a cross-functional engineering team to unify business operations from three distinct divisions onto a singular platform, developing 60+ RESTful services and utilizing MuleSoft to securely migrate 900,000+ data artifacts.
- Architected zero-downtime deployment strategies and robust fault-tolerant pipelines during a massive data migration project encompassing roughly 65 million individual file assets.

Software Integration Engineer - Synergis Technologies

March 2012 - July 2016

- Optimized data journaling and big data analytics utilizing Hadoop and MapReduce frameworks, successfully compressing data processing cycles from 2.5 hours down to 12 minutes - a 92% reduction in processing time.
 - Modernized roughly 60 internal applications from legacy procedural languages to scalable, cloud-ready frameworks, drastically improving system performance and concurrent user capacity.
-

EDUCATION

Master of Business Administration (MBA) - University of Washington

2021

Master of Science in Computer Science - Georgia Institute of Technology

2011

CERTIFICATIONS

- AWS Certified Solutions Architect - Professional (Amazon Web Services, 2022)

PROJECTS

Algorithmic Trading Agent

Built an autonomous real-time market analysis agent capable of identifying historic price action patterns, generating predictive models, and proposing systematic trade entries.

Proprietary EVAL Framework

Engineered a rigorous AI evaluation and benchmarking system to score emerging LLMs against internal safety, hallucination, and accuracy thresholds prior to production deployment.

Custom MCP Integration

Developed a proprietary Model Context Protocol (MCP) server enabling specialized AI agents to securely query real-time financial APIs and generate dynamic charting and data visualizations for market analysis.

Multi-Modal Content Moderation

Created an automated content moderation pipeline leveraging Google's multimodal analysis models, integrated into an agentic workflow to autonomously review and flag video and image assets.

Enterprise Copilot Scaling

Drove the organizational rollout of GitHub Copilot to over 500 developers by engineering a centralized governance framework, enforcing VS Code guardrails and automated compliance checks.

Secure PKI Chat Network

Architected a decentralized web platform utilizing Public Key Infrastructure (PKI) for cryptographic token-based member verification, securing private enterprise communications.